

Drop-in Guide

TRACK	ROBOT	LAYER	DEADLINE
● agents	unitree go2	interaction	may 28 · 18:00

A calm dispatch terminal for an agentic guide robot. Drop a robot in any building. It learns the space by asking questions. It guides any visitor through it — and waits if you fall behind.

"It's not a delivery bot. It's a guide that knows when to wait for you."

DROP-IN GUIDE · V1.0 · FIELD-TESTED AT VENUE

The robot authors its own scene memory.

TODAY

system integrators

Existing embodied-AI guides (AGIBOT, hospital wayfinders, museum robots) need pre-mapped environments and operator-authored knowledge bases. The deployment model is "send a system integrator, spend a week."

DROP-IN GUIDE

operator + AI together


The robot proposes tags from a 5-minute walkthrough; the operator confirms with a keypress. Same principle as Rehnova's operator-grounded verification – but in physical space, not spreadsheet rows.

Where this *project* lives.

MOTION	Unitree Go2 hardware + dimOS nav stack — <i>A* planning, frontier exploration, PG0 loop closure. Locomotion, obstacle avoidance, low-level control. Provided.</i>
TASK	dimOS skill runtime + Claude Sonnet 4.6 via MCP — <i>decides which skill to call when, manages dialogue state, parses natural language into tool sequences. Provided.</i>
INTERACTION ★	This PR. <i>Generative priming dialogue, grounded narration, decision audit, voice-triggered "I'll wait for you", auto-arrival greeting. The interaction layer is where we add value.</i>

Priming. Guidance. Q&A.

The same *Claude system prompt* drives all three phases — text, voice, and WASD keypress inputs produce identical agent behavior, including the calibrated-uncertainty checks and the *[SYSTEM ARRIVAL EVENT]* chain.



The pause is what makes it a guide.

A delivery bot drives to a destination. A guide notices when you've fallen behind. We can't see backward through the Go2's front camera, so the visitor controls the pace with their voice — the *natural way a human guide would work*.

```
# validated end-to-end in dimOS replay
visitor: take me to the copier
robot:   log_nav_decision(copier, semantic, 0.4)
        speak("Making my best guess — follow me!")
        navigate_with_text("copier")
        [walking...]

visitor: wait, hold on, I'm behind
robot:   stop_navigation()
        ► speak("I'll wait for you here.") ← signature line

visitor: okay, I'm ready, let's go
robot:   log_nav_decision(copier, semantic, 0.4)
        navigate_with_text("copier") ← resumes same target

[planner] /goal_reached = True
```

One blueprint. Six skills. 1,330 lines.

BLUEPRINT	drop_in_guide.py — composes from unitree_go2 base, skips CUDA-only modules, registers 3-phase system prompt.
SKILLS	scene_caption · teleop_velocity · arrival_announcer · decision_audit · lead_with_follow · reactive_qa
VALIDATED	Replay run on data/go2_short.db: 22 modules load, /goal_reached chain fires, voice-wait + resume both verified. No hallucinated arrival in any test.
FIELD	Tested live on Go2 at venue. obstacle_avoidance=False baked into config. macOS multicast workaround documented.

PR · DIMENSIONALOS/DIMOS #2289

TOOLING · AROME3/DROP-IN-GUIDE